Session 2: Safe Harbor policy sprint

AI Red-Teaming, Governance, Ethics & Privacy

Your outcome

Use three realistic scenarios to draft a one-page organizational AI policy, then test whether a colleague can follow it.

Before you begin

Bring your Session 1 use case and your organization's existing AI/data rules, if available. Use descriptions of data categories rather than actual private records.

Participant details

Name: _____ Cohort: _____
Project / task: _____

Working boundaries

Use fictional or explicitly approved examples. Do not paste private customer, employee, health or financial records into public AI tools. Follow your organization's tool and data rules. Treat outputs as drafts; a person verifies facts and approves external actions.

What to keep

Keep a copy of your inputs, outputs, decisions and revisions. Browser-only drafts may not follow you to another device. Upload the completed resource or share a permitted link through your participant dashboard.

Course coaching

Each cohort participant may request one remote coaching session during the course. Request it inside the platform; your instructor confirms the time. Use the same conversation for time changes.

Lab guide

Follow your instructor or work through these steps using approved tools. A document or manual simulation is acceptable when a connected tool is unavailable; label its status honestly.

Step 1

List the tools, tasks, and data categories the policy should cover. Separate what the assistant may draft from what it may send, publish, or change.

Step 2

For any connected tool, name the account owner, minimum required access, approval point, and way to revoke access. Write allowed, prohibited, and ask-first examples.

Step 3

Test the draft against a routine task, a sensitive-data request, and an uncertain case. A reusable playbook does not grant permission to act.

Worked example — fictional / illustrative

Fictional example: Northstar permits brainstorming with invented examples in approved tools, prohibits uploading customer records to public tools, and requires the operations lead to review exceptions. This is an exercise, not legal advice or organizational approval.

Try the prompt sheet

Use the companion prompt sheet as a starting point, replace bracketed fields and test the result. It is not a guarantee of accuracy or an instruction to publish anything.

Participant worksheet

Complete the fields below. Use the editable Word version if you need more space. Label evidence you have not collected yet as unknown.

Policy owner / version / approval status

_____
_____

Scope: people, tasks and tools covered

_____
_____

Approved tools and who authorizes access

_____
_____

Permitted data / prohibited data / uncertain cases

_____
_____

Human checks and accountable approver

_____
_____

Participant worksheet — continued

Complete the fields below. Use the editable Word version if you need more space. Label evidence you have not collected yet as unknown.

Disclosure, copyright and accessibility expectations

_____
_____

Incident: stop, preserve relevant evidence, notify whom

_____
_____

Exception request and escalation route

_____
_____

Peer scenario results / revision / approval still needed

_____
_____

Test, challenge and improve

Test 1

A colleague wants to summarize a customer complaint containing personal information in a public chatbot. What must be removed, approved or avoided?

Expected result: _____
Actual result / evidence: _____
Correction or decision: _____

Test 2

A marketing draft includes a convincing but unsourced customer quote. Decide who verifies it and what happens if it cannot be verified.

Expected result: _____
Actual result / evidence: _____
Correction or decision: _____

Test 3

An employee asks AI to rank job applicants from their CVs. Identify the employment, privacy and fairness risks and the escalation owner before proceeding.

Expected result: _____
Actual result / evidence: _____
Correction or decision: _____

Transfer challenge

Ask a peer to decide what to do in one scenario using only your policy. Replace ambiguous wording before seeking organizational approval.

Review and submit

Mark each criterion ready, needs revision, or not applicable with a reason. An instructor reviews the evidence; completing a worksheet alone does not establish approval or certificate eligibility.

[ ] Ready [ ] Revise [ ] N/A — Allowed tools and use cases are named or explicitly awaiting approval.

[ ] Ready [ ] Revise [ ] N/A — Prohibited data and actions have concrete examples.

[ ] Ready [ ] Revise [ ] N/A — Human review and escalation ownership are clear.

[ ] Ready [ ] Revise [ ] N/A — Three scenarios demonstrate how the policy is applied.

[ ] Ready [ ] Revise [ ] N/A — The policy's approval status and next review date are recorded.

Submission

Submit under Safe Harbor Usage Policy in My work. Include your artifact or permitted link, test evidence and notes about what changed. Save a draft if it is not ready for review.

Feedback and follow-through

Feedback received: _____
Change I made: _____
Evidence retested: _____
My next action / owner / review point: _____